

```
In [41]: import numpy as np
        from tensorflow import keras
        from tensorflow.keras import layers

In [42]: # Set the number of categories and define the shape of input images
        num_categories = 10
        input_shape = (28, 28, 1)

In [43]: # Load the MNIST dataset and split it into training and testing sets
        (train_images, train_labels), (test_images, test_labels) = keras.datasets.mnist.load_data()

In [44]: # Normalize the pixel values of the images to a range between 0 and 1
        train_images = train_images.astype("float32") / 255.0
        test_images = test_images.astype("float32") / 255.0

In [45]: # Expand the dimensions of the images to include the color channel
        train_images = np.expand_dims(train_images, -1)
        test_images = np.expand_dims(test_images, -1)

In [46]: # Convert class labels to categorical matrices
        train_labels = keras.utils.to_categorical(train_labels, num_categories)
        test_labels = keras.utils.to_categorical(test_labels, num_categories)

In [47]: # Define a modified convolutional neural network architecture
        model = keras.Sequential([
            keras.Input(shape=input_shape),
            layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
            layers.MaxPooling2D(pool_size=(2, 2)),
            layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
            layers.MaxPooling2D(pool_size=(2, 2)),
            layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
            layers.Flatten(),
            layers.Dropout(0.5),
            layers.Dense(256, activation="relu"),
            layers.Dense(num_categories, activation="softmax")
        ])

In [48]: # Display the summary of the model architecture
        model.summary()

Model: "sequential_2"

Layer (type)                Output Shape                Param #
=====
conv2d_6 (Conv2D)           (None, 26, 26, 32)         320

max_pooling2d_4 (MaxPoolin  (None, 13, 13, 32)         0
g2D)

conv2d_7 (Conv2D)           (None, 11, 11, 64)         18496

max_pooling2d_5 (MaxPoolin  (None, 5, 5, 64)           0
g2D)

conv2d_8 (Conv2D)           (None, 3, 3, 128)          73856

flatten_2 (Flatten)         (None, 1152)                0

dropout_2 (Dropout)         (None, 1152)                0

dense_4 (Dense)             (None, 256)                 295168

dense_5 (Dense)             (None, 10)                  2570
=====
Total params: 390410 (1.49 MB)
Trainable params: 390410 (1.49 MB)
Non-trainable params: 0 (0.00 Byte)

In [49]: # Compile the model with specified optimizer and loss function
        model.compile(optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"])

In [50]: # Set the batch size and number of epochs for training
        batch_size = 128
        epochs = 15

In [51]: # Train the model using the training data
        model.fit(train_images, train_labels, batch_size=batch_size, epochs=epochs)

Epoch 1/15
469/469 [=====] - 20s 41ms/step - loss: 0.2338 - accuracy: 0.9268
Epoch 2/15
469/469 [=====] - 19s 40ms/step - loss: 0.0647 - accuracy: 0.9803
Epoch 3/15
469/469 [=====] - 19s 40ms/step - loss: 0.0501 - accuracy: 0.9840
Epoch 4/15
469/469 [=====] - 20s 43ms/step - loss: 0.0431 - accuracy: 0.9866
Epoch 5/15
469/469 [=====] - 20s 43ms/step - loss: 0.0342 - accuracy: 0.9900
Epoch 6/15
469/469 [=====] - 20s 42ms/step - loss: 0.0304 - accuracy: 0.9904
Epoch 7/15
469/469 [=====] - 20s 42ms/step - loss: 0.0272 - accuracy: 0.9915
Epoch 8/15
469/469 [=====] - 20s 42ms/step - loss: 0.0228 - accuracy: 0.9926
Epoch 9/15
469/469 [=====] - 20s 42ms/step - loss: 0.0224 - accuracy: 0.9929
Epoch 10/15
469/469 [=====] - 19s 40ms/step - loss: 0.0198 - accuracy: 0.9934
Epoch 11/15
469/469 [=====] - 19s 41ms/step - loss: 0.0172 - accuracy: 0.9943
Epoch 12/15
469/469 [=====] - 20s 42ms/step - loss: 0.0164 - accuracy: 0.9945
Epoch 13/15
469/469 [=====] - 20s 42ms/step - loss: 0.0144 - accuracy: 0.9951
Epoch 14/15
469/469 [=====] - 22s 47ms/step - loss: 0.0144 - accuracy: 0.9950
Epoch 15/15
469/469 [=====] - 20s 42ms/step - loss: 0.0122 - accuracy: 0.9963
Out[51]: <keras.src.callbacks.History at 0x1e7cd216390>

In [52]: # Evaluate the model performance on the test data
        evaluation = model.evaluate(test_images, test_labels, verbose=0)
        print("Test loss:", evaluation[0])
        print("Test accuracy:", evaluation[1])

Test loss: 0.021707022562623024
Test accuracy: 0.9937000274658203

In [53]: # Save the trained model for future use
        model.save('D:\PRABIN CNN')

INFO:tensorflow:Assets written to: D:\PRABIN CNN\assets
INFO:tensorflow:Assets written to: D:\PRABIN CNN\assets

In [54]: model.save('MNIST_baseline_model.h5')

C:\Users\prabi\anaconda3\Lib\site-packages\keras\src\engine\training.py:3079: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
  saving_api.save_model(

In [55]: model_path = 'D:\Hand-digit\my_model.h5' # Replace with your desired path

In [56]: model.save(model_path)

In [57]: from tensorflow.keras.models import load_model

        model = load_model('D:\Hand-digit\my_model.h5') # Replace with your model file

In [58]: from tensorflow.keras.preprocessing.image import img_to_array, load_img
        import numpy as np

        img = load_img('C:/Users/prabi/Downloads/sample_digit_7.png', color_mode='grayscale', target_size=(28, 28))

        # Convert the image to an array and normalize
        img_array = img_to_array(img) / 255.0

        # Reshape the array for the model
        img_array = img_array.reshape((1, 28, 28, 1)) # The first 1 is for the batch size

In [59]: prediction = model.predict(img_array)

1/1 [=====] - 0s 58ms/step

In [60]: predicted_class = np.argmax(prediction, axis=1)

In [61]: print(f"The model predicts: {predicted_class[0]}")

The model predicts: 7

In [ ]:
```